

# Technical Document: Lab 1

## Overview

### Team

Role:	Name:
Manager	Luis Arcos
Technical Lead	Madylin Teel
Developer	Coleman Bixler
Developer	Garrett Snow
Developer	An Vu
Developer	Nathan Stephani

## Conventions

### Naming

Branch names	Example: {dev,techlead,manager}.<user>/<objective>/<task> i.e., dev.madyteel/obj-a/t-01
Generated output folders	Example: task##/<operation>-versions/ i.e., task02/split-versions/
C code file names	Example: var##-<operation/schedule>.x i.e., var00-ssa-schedule-blur3x3.c

## Objective A: Handspun Simple Schedule Code Generator

Objective A focuses on the process of applying schedule transformation (by hand) on an operation in Static Single Assignment (SSA) form.

**Static Single Assignment (SSA)** – A representation used in compilers where each variable is assigned exactly once and defined before it is used.

Please see `LAB-SCHEDULE-DECOUPLED/lab\_initial/operations/` for all operations which need to be converted to SSA format.

Please see `LAB-SCHEDULE-DECOUPLED/lab\_initial/examples/basic-ssa/{ssa\_var000.c, ssa\_var001.c}` for examples on how to convert the operations into SSA form.

## T00: Handspun Baseline

Implement, Verify, Benchmark, and Compare for All of the Operation Specifications in SSA form.

### Implementation

Use `ssa\_var000.c` and `ssa\_var001.c` as references for structuring loop transformations while maintaining SSA-compliance [from lab\_initial/examples/basic-ssa/].

#### Steps:

1. Convert each individual operation into SSA form:
  - a. Transform each operation file into SSA-compliant C code.
    - i. Rewrite each computation such that every value is assigned exactly and only once.
    - ii. Introduce new variables instead of reassigning existing ones.
  - b. Ensure there is no variable reassignment!
2. Implement the SSA-compliant C code:
  - a. Must write manual implementations of the SSA transformations.
  - b. Can verify against `baseline.c` [from lab\_initial/examples/basic-ssa/].
  - c. Store all ssa versions in their own folder [i.e., task00/ssa-versions/].
  - d. Ensure consistent naming across versions [i.e., `ssa-var00`].
3. Verify correctness:
  - a. Run `run-all-verify.sh`. (WIP)
4. Benchmark performance:
  - a. Run `run-all-measure.sh`. (WIP)
5. Compare SSA implementations:
  - a. Compare the performance across various operations (such as blur, conv, jacobi2d, etc.).
  - b. Identify which operations benefit the most from SSA transformation.
6. Clean up generated files:
  - a. Run `run-clean-all.sh` to clean up all files generated by a given area. (WIP)

### Benchmarking

## T01: Handspun + Unroll

Implement, Verify, Benchmark, and Compare for Five (5) of the Operation Specifications in SSA form.

### Implementation

#### Steps:

1. If not already done, convert each operation into SSA form (see objective-a T00 for more details):

- a. Each computation must be assigned to a unique variable.
  - b. No variable reassignment is allowed!
2. Apply loop unrolling:
 

Unrolling reduces loop overhead by manually duplicating loop iterations.

  - a. Select any (5) operations [from /lab\_initial/operations/]. (For example: blur, conv, jacobi2d, etc.)
  - b. Manually unroll the inner loop:
    - i. Identify which loop needs to be unrolled (either j0, q0, or r0 for example from slides)
    - ii. Example given in the slides for this lab, specifically slide #6.

## Schedules: unroll <id>

The following will take the code on the left and transform it to the code on the right. Loops will be labelled by their index variable.

unroll r0

```
void COMPUTE_NAME( int m0,
                  int n0,
                  float *x,
                  float *y )
{
    // Q=3xR=4
    float weights[] =
    //r=0      1      2      3
    {-1.1, -1.1, 1.2, -2.1, // q=0
     1.1, -2.1, -1.2, 2.2, // q=1
     2.1, 0.1, 0.2, 1.2}; // q=2

    BEGIN_INSTRUMENTATION;
    for( int i0 = 0; i0 < m0; ++i0 )
        for( int j0 = 0; j0 < n0; ++j0 )
            for( int q0 = 0; q0 < (Q); ++q0 )
                for( int r0 = 0; r0 < (R); ++r0 )
                    y[i0*n0+j0] += weights[q0*(R)+r0] *
                        x[(q0*i0)%m0 + ((r0+j0)%n0) ];
    END_INSTRUMENTATION;
}
```

```
void COMPUTE_NAME( int m0,
                  int n0,
                  float *x,
                  float *y )
{
    // BEGIN_INSTRUMENTATION; // func: COMPUTE_NAME
    for( int i0 = 0; i0 < m0; ++i0 )
        // BEGIN_INSTRUMENTATION; // loop:i0
        for( int j0 = 0; j0 < n0; ++j0 )
            // BEGIN_INSTRUMENTATION; // loop:j0
            for( int q0 = 0; q0 < (Q); ++q0 )
                {
                    BEGIN_INSTRUMENTATION; // loop:q0
                    {
                        // r0=0
                        y[i0*n0+j0] += weights[q0*(R)+0] *
                            x[(q0*i0)%m0 + ((0+j0)%n0) ];
                    }
                    // r0=1
                    y[i0*n0+j0] += weights[q0*(R)+1] *
                        x[(q0*i0)%m0 + ((1+j0)%n0) ];
                    // r0=2
                    y[i0*n0+j0] += weights[q0*(R)+2] *
                        x[(q0*i0)%m0 + ((2+j0)%n0) ];
                    // r0=3
                    y[i0*n0+j0] += weights[q0*(R)+3] *
                        x[(q0*i0)%m0 + ((3+j0)%n0) ];
                }
            END_INSTRUMENTATION; // loop:q0
        // END_INSTRUMENTATION; // loop:j0
    // END_INSTRUMENTATION; // loop:i0
    // END_INSTRUMENTATION; // func: COMPUTE_NAME
}
```

3. Implement the SSA-compliant unrolled C code:
  - i. Store all unrolled implementations in their own folder [i.e., /task01/unrolled-versions/].
  - ii. Ensure consistent naming across versions [i.e., `unroll-var00`].
4. Verify correctness:
  - a. Run `run-all-verify.sh`. (WIP)
  - b. Ensure unrolling does not introduce any unexpected behavior.
5. Benchmark performance:
  - a. Run `run-all-measure.sh`. (WIP)
    - i. Unrolled SSA vs. non-unrolled SSAX
    - ii. Unrolled SSA vs. Baseline (non-SSA)
6. Compare results:
  - a. Compare the execution time, FLOPs, and memory usage for:
    - i. Unrolled SSA vs. non-unrolled SSA
    - ii. Unrolled SSA vs. Baseline (non-SSA)
    - iii. Which operations benefit the most from loop unrolling?
    - iv. Are there any cases in which unrolling hurts performance?
7. Clean up generated files:
  - a. Run `run-clean-all.sh` to clean up all files generated by a given area. (WIP)

## T02: Handspun + Split

Implement, Verify, Benchmark, and Compare for Five (5) of the Operation Specifications in SSA form.

### Implementation

#### Steps:

1. If not already done, convert each operation into SSA form (see objective-a T00 for more details):
  - a. Each computation must be assigned to a unique variable.
  - b. No variable reassignment is allowed!
2. Apply loop splitting:
 

Splitting reduces iteration overhead and improves cache locality by breaking a single loop into two nested loops.

  - a. Select any (5) operations [from /lab\_initial/operations/]. (For example: blur, conv, jacobi2d, etc.)
  - b. Manually split the loops:
    - i. Identify which loops need to be split (either j0, q0, or r0 for example from slides)
    - ii. Example given in the slides for this lab, specifically slide #7.

### Schedules: split <id> <out-id> <inner-id> <factor>

The following will take the code on the left and transform it to the code on the right. Loops will be labelled by their index variable.

split j0 j0\_o j0\_i 2

```

void COMPUTE_NAME( int m0,
                   int n0,
                   float *x,
                   float *y )
{
    // Q=3xR=4
    float weights[] =
    //r=0  1  2  3
    {-1.1, -1.1, 1.2, -2.1, // q=0
     1.1, 2.1, 1.2, 2.2, // q=1
     2.1, 0.1, 0.2, 1.2}; // q=2

    BEGIN_INSTRUMENTATION;
    for( int i0 = 0; i0 < m0; ++i0 )
        for( int j0 = 0; j0 < n0; ++j0 )
            for( int q0 = 0; q0 < (Q); ++q0 )
                for( int r0 = 0; r0 < (R); ++r0 )
                    y[i0*n0+j0] += weights[q0*(R)+r0] *
                    x[ ((q0+i0)%m0)*n0 + ((r0+j0)%n0) ];
    END_INSTRUMENTATION;
}

```

```

void COMPUTE_NAME( int m0,
                   int n0,
                   float *x,
                   float *y )
{
    for( int i0 = 0; i0 < m0; ++i0 )
    {
        for( int j0_o = 0; j0_o < n0; j0_o+=2 )
            for( int j0_i = 0; j0_i < 2; ++j0_i )
            {
                // BEGIN_INSTRUMENTATION; // loop:j0
                for( int q0 = 0; q0 < (Q); ++q0 )
                {
                    // BEGIN_INSTRUMENTATION; // loop:q0
                    for( int r0 = 0; r0 < (R); ++r0 )
                    {
                        BEGIN_INSTRUMENTATION; // loop:r0
                        y[i0*n0+(j0_o+j0_i)] += weights[q0*(R)+r0] *
                        x[ ((q0+i0)%m0)*n0 + ((r0+(j0_o+j0_i))%n0) ];
                        END_INSTRUMENTATION; // loop:r0
                    }
                }
            }
    }
}

```

3. Implement the SSA-compliant split C code:
  - i. Store all split implementations in their own folder [i.e., /task02/split-versions/].
  - ii. Ensure consistent naming across versions [i.e., `split-var00`].
4. Verify correctness:
  - a. Run `run-all-verify.sh`. (WIP)
  - b. Ensure loop splitting does not introduce any unexpected behavior.
5. Benchmark performance:
  - a. Run `run-all-measure.sh`. (WIP)

- i. Split SSA vs. non-split SSA
  - ii. Split SSA vs. Baseline (non-SSA)
- 6. Compare results:
  - a. Compare the execution time, FLOPs, and memory usage for:
    - i. Split SSA vs. non-split SSA
    - ii. Split SSA vs. Baseline (non-SSA)
    - iii. Which operations benefit the most from loop splitting?
    - iv. Are there any cases in which splitting hurts performance perhaps by introducing overhead?
- 7. Clean up generated files:
  - a. Run `run-clean-all.sh` to clean up all files generated by a given area. (WIP)

## T03: Handspun + Interchange

Implement, Verify, Benchmark, and Compare for One (1) Operation against six (6) different ways to interchange the loops in SSA form.

## Implementation

### Steps:

1. If not already done, convert the operation into SSA form (see objective-a T00 for more details):
  - a. Each computation must be assigned to a unique variable.
  - b. No variable reassignment is allowed!
2. Apply loop interchange in (6) different ways:
 

Loop interchange affects both cache locality and execution efficiency by changing the order of nested loops.

  - a. Select any (1) operation [from /lab\_initial/operations/]. (For example: blur, conv, jacobi2d, etc.)
  - b. Manually interchange the loops:
    - i. Identify which loops need to be interchanged.
    - ii. Example given in the slides for this lab, specifically slide #8.

### Schedules: interchange <id-a> <id-b>

The following will take the code on the left and transform it to the code on the right. Loops will be labelled by their index variable.

interchange r0 j0

```
void COMPUTE_NAME( int m0,
                  int n0,
                  float *x,
                  float *y )
{
    // Q=3xR=4
    float weights[] =
    {
        // r0
        1, 2, 3
    };
    [-1.1, -1.1, 1.2, -2.1, // q0
     -1.1, -2.1, 1.2, 2.2, // q1
     -2.1, 0.1, 0.2, 1.2]; // q2

    BEGIN_INSTRUMENTATION;
    for( int i0 = 0; i0 < m0; ++i0 )
        for( int j0 = 0; j0 < n0; ++j0 )
            for( int q0 = 0; q0 < Q; ++q0 )
                for( int r0 = 0; r0 < R; ++r0 )
                    y[(i0*n0+j0) + weights[q0*(R)+r0] *
                     x[(q0*10+m0)*n0 + ((r0*j0)%n0) ];
    END_INSTRUMENTATION;
}
```

```
void COMPUTE_NAME( int m0,
                  int n0,
                  float *x,
                  float *y )
{
    // BEGIN_INSTRUMENTATION; // func: COMPUTE_NAME
    for( int i0 = 0; i0 < m0; ++i0 )
    {
        // BEGIN_INSTRUMENTATION; // loop:i0
        for( int r0 = 0; r0 < R; ++r0 )
        {
            // BEGIN_INSTRUMENTATION; // loop:r0
            for( int q0 = 0; q0 < Q; ++q0 )
            {
                // BEGIN_INSTRUMENTATION; // loop:q0
                for( int j0 = 0; j0 < n0; ++j0 )
                {
                    BEGIN_INSTRUMENTATION; // loop:j0
                    y[(i0*n0+j0) + weights[q0*(R)+r0] *
                     x[(q0*10+m0)*n0 + ((r0*j0)%n0) ];
                    END_INSTRUMENTATION; // loop:j0
                }
            }
            // END_INSTRUMENTATION; // loop:q0
        }
        // END_INSTRUMENTATION; // loop:r0
    }
    // END_INSTRUMENTATION; // loop:i0
    // END_INSTRUMENTATION; // func: COMPUTE_NAME
}
```

3. Implement the SSA-compliant split C code:
  - i. Store all interchange implementations in their own folder [i.e., /task03/interchange-versions/].
  - ii. Ensure consistent naming across versions [i.e., `ssa-interchange-var00`].
4. Verify correctness:
  - a. Run `run-all-verify.sh`. (WIP)
  - b. Ensure loop interchanges do not introduce any unexpected behavior.
5. Benchmark performance:
  - a. Run `run-all-measure.sh`. (WIP)
    - i. Each interchanged version vs. baseline SSA
    - ii. Each interchanged version vs. other interchanged versions
6. Compare results:
  - a. Compare the execution time, FLOPs, and memory usage for:
    - i. Each interchanged version vs. baseline SSA
    - ii. Which interchanges give the most optimal results?
    - iii. Are there any cases in which interchanging worsens performance?
7. Clean up generated files:
  - a. Run `run-clean-all.sh` to clean up all files generated by a given area. (WIP)

## T04: Handspun + Complex Schedule (3 Commands)

Implement, Verify, Benchmark, and Compare for three (3) of the Operation Specifications using the same schedule in SSA form.

### Implementation

#### Steps:

1. Before applying complex scheduling, ensure the (3) desired operations are in SSA form (see objective-a T00 for more details):
  - a. Each computation must be assigned to a unique variable.
  - b. No variable reassignment is allowed!
2. Either define or utilize premade complex schedule:
 

A complex schedule applies (in this case, 3) transformations to the loop structure (which could be any combination of: unrolling, splitting, interchanging, etc.).

Predefined schedules can be found in /lab\_initial/schedules/complex\_mix/. *The first two options in this folder are (3) command schedules.*

Example of a (3) command complex schedule:

- a. Apply loop splitting: split j0 into j0\_outer and j0\_inner
  - b. Apply loop interchange: swap q0 and j0\_outer
  - c. Apply loop unrolling: unroll j0\_inner by a factor of 2
3. Apply the complex schedule to (3) operations:
  - a. Select any (3) operations [from /lab\_initial/operations/]. (For example: blur, conv, jacobi2d, etc.)
  - b. Apply the same transformation sequence across all (3) operations.

- c. Example given in the slides for this lab, specifically slide #9 and slide #10.

## Schedules: Multiple Commands

Multiple scheduling commands can be used and are **executed from top to bottom**.

**split j0 j0\_o j0\_i 2**

**interchange r0 j0\_i**

**unroll j0\_i**

```
void COMPUTE_NAME( int m0,
                   int n0,
                   float *x,
                   float *y )
{
    // Q=3xR=4
    float weights[] =
    {
        // r=0
        1, 2, 3
        {-1.1, -1.1, 1.2, -2.1, // q=0
         -1.1, -2.1, 1.2, 2.2, // q=1
         -2.1, 0.1, 0.2, 1.2}; // q=2
    };

    BEGIN_INSTRUMENTATION;
    for( int i0 = 0; i0 < m0; ++i0 )
        for( int j0 = 0; j0 < n0; ++j0 )
            for( int q0 = 0; q0 < (Q); ++q0 )
                for( int r0 = 0; r0 < (R); ++r0 )
                    y[i0*n0+j0] += weights[q0*(R)+r0] *
                        x[(q0-i0)%m0*n0 + ((r0+j0)%n0)];
    END_INSTRUMENTATION;
}
```

```
void COMPUTE_NAME( int m0,
                   int n0,
                   float *x,
                   float *y )
{
    for( int i0 = 0; i0 < m0; ++i0 )
    {
        for( int j0_o = 0; j0_o < n0; j0_o+=2 )
        {
            for( int r0 = 0; r0 < (R); ++r0 )
            {
                for( int q0 = 0; q0 < (Q); ++q0 )
                {
                    BEGIN_INSTRUMENTATION; // loop:q0
                    { // j0_i=0
                        y[i0*n0+(j0_o+0)] += weights[q0*(R)+r0] *
                            x[(q0-i0)%m0*n0 + ((r0+(j0_o+0))%n0)];
                    }
                    { // j0_i=1
                        y[i0*n0+(j0_o+1)] += weights[q0*(R)+r0] *
                            x[(q0-i0)%m0*n0 + ((r0+(j0_o+1))%n0)];
                    }
                }
            }
        }
        END_INSTRUMENTATION; // loop:c0
    }
}
```

## Schedules: Multiple Commands

conv

```
void COMPUTE_NAME( int m0,
                   int n0,
                   float *x,
                   float *y )
{
    // Q=3xR=4
    float weights[] =
    {
        // r=0
        1, 2, 3
        {-1.1, -1.1, 1.2, -2.1, // q=0
         -1.1, -2.1, 1.2, 2.2, // q=1
         -2.1, 0.1, 0.2, 1.2}; // q=2
    };

    BEGIN_INSTRUMENTATION;
    for( int i0 = 0; i0 < m0; ++i0 )
        for( int j0_o = 0; j0_o < n0; ++j0_o )
            for( int q0 = 0; q0 < (Q); ++q0 )
                for( int r0 = 0; r0 < (R); ++r0 )
                    y[i0*n0+j0_o] += weights[q0*(R)+r0] *
                        x[(q0-i0)%m0*n0 + ((r0+j0_o)%n0)];
    END_INSTRUMENTATION;
}
```

**split j0 j0\_o j0\_i 2**

**interchange r0 j0\_i**

**unroll j0\_i**

opt1 = Split(j0,j0\_o,j0\_i,2, conv)

```
void COMPUTE_NAME( int m0,
                   int n0,
                   float *x,
                   float *y )
{
    for( int i0 = 0; i0 < m0; ++i0 )
    {
        for( int j0_o = 0; j0_o < n0; j0_o+=2 )
        {
            for( int r0 = 0; r0 < (R); ++r0 )
            {
                for( int q0 = 0; q0 < (Q); ++q0 )
                {
                    BEGIN_INSTRUMENTATION; // loop:q0
                    { // j0_i=0
                        y[i0*n0+(j0_o+0)] += weights[q0*(R)+r0] *
                            x[(q0-i0)%m0*n0 + ((r0+(j0_o+0))%n0)];
                    }
                    { // j0_i=1
                        y[i0*n0+(j0_o+1)] += weights[q0*(R)+r0] *
                            x[(q0-i0)%m0*n0 + ((r0+(j0_o+1))%n0)];
                    }
                }
            }
        }
    }
}
```

opt2 = Interchange(r0,j0\_i, opt1)

```
void COMPUTE_NAME( int m0,
                   int n0,
                   float *x,
                   float *y )
{
    for( int i0 = 0; i0 < m0; ++i0 )
    {
        for( int j0_o = 0; j0_o < n0; j0_o+=2 )
        {
            for( int r0 = 0; r0 < (R); ++r0 )
            {
                for( int q0 = 0; q0 < (Q); ++q0 )
                {
                    BEGIN_INSTRUMENTATION; // loop:q0
                    { // j0_i=0
                        y[i0*n0+(j0_o+0)] += weights[q0*(R)+r0] *
                            x[(q0-i0)%m0*n0 + ((r0+(j0_o+0))%n0)];
                    }
                    { // j0_i=1
                        y[i0*n0+(j0_o+1)] += weights[q0*(R)+r0] *
                            x[(q0-i0)%m0*n0 + ((r0+(j0_o+1))%n0)];
                    }
                }
            }
        }
    }
}
```

opt3 = Unroll(j0\_i, opt2)

```
void COMPUTE_NAME( int m0,
                   int n0,
                   float *x,
                   float *y )
{
    for( int i0 = 0; i0 < m0; ++i0 )
    {
        for( int j0_o = 0; j0_o < n0; j0_o+=2 )
        {
            for( int r0 = 0; r0 < (R); ++r0 )
            {
                for( int q0 = 0; q0 < (Q); ++q0 )
                {
                    BEGIN_INSTRUMENTATION; // loop:q0
                    { // j0_i=0
                        y[i0*n0+(j0_o+0)] += weights[q0*(R)+r0] *
                            x[(q0-i0)%m0*n0 + ((r0+(j0_o+0))%n0)];
                    }
                    { // j0_i=1
                        y[i0*n0+(j0_o+1)] += weights[q0*(R)+r0] *
                            x[(q0-i0)%m0*n0 + ((r0+(j0_o+1))%n0)];
                    }
                }
            }
        }
    }
}
```

4. Implement SSA-compliant scheduled code:
  - a. Store all scheduled implementations in their own folder [i.e., /task04/scheduled-versions/].
  - b. Ensure consistent naming across versions [i.e., `ssa-schedule-blur3x3.c`].
5. Verify correctness:
  - a. Run `run-all-verify.sh`. (WIP)
  - b. Ensure complex schedule does not introduce any unexpected behavior.
6. Benchmark performance:
  - a. Run `run-all-measure.sh`. (WIP)



- i. Scheduled SSA vs. non-scheduled SSA
  - ii. Scheduled SSA vs. baseline (non-SSA)
- 7. Compare results:
  - a. Compare the execution time, FLOPs, and memory usage for:
    - i. Scheduled SSA vs. non-scheduled SSA
    - ii. Scheduled SSA vs. baseline (non-SSA)
    - iii. Which combinations of transformations performs the best?
    - iv. Are there any cases in which a specific operation benefits more from the schedule?
- 8. Clean up generated files:
  - a. Run ``run-clean-all.sh`` to clean up all files generated by a given area. (WIP)

## Optional: For Additional Points

### *T01-Opt00: Handspun + Unroll(r) vs Unroll(q) [Optional]*

Implement, Verify, Benchmark, and Compare for One (1) Operation, but two (2) different loops in SSA form.

### *T02-Opt00: Handspun + Split(i) vs Split(j) vs Split(q) vs Split(r) [Optional]*

Implement, Verify, Benchmark, and Compare for One (1) Operation, but four (4) different loops to be split in SSA form.

### *T02-Opt01: Handspun + Split(2) vs Split(4) vs Split(8) [Optional]*

Implement, Verify, Benchmark, and Compare for One (1) Operation, One (1) loop, but four (4) different split factors) in SSA form.

### *T04-Opt00: Handspun + Complex Schedule (3 Commands) [Optional]*

Implement, Verify, Benchmark, and Compare for one (1) of the Operation Specification using three (3) different schedules in SSA form.

## Objective B: Standalone Simple Schedule Code Generator

The goal of this objective is to focus on automating what was previously done manually in objective-a.

### T00: Simple Standalone Baseline

Implement, Verify, Benchmark, and Compare for All of the Operation Specifications in SSA form.



# Objective B: Standalone Simple Schedule

## Operation Specification

```

3 4
-1.1 -1.1 1.2 -2.1
-1.1 -2.1 -1.2 2.2
-2.1 0.1 0.2 1.2
i0 j0 q0 r0

```

Standalone Code  
Generator

```

void COMPUTE_NAME( int m0,
                  int n0,
                  float *x,
                  float *y )
{
    // Q=3xR-4
    float weights[] =
    //r=0  1  2  3
    { 1.1, -1.1, 1.2, -2.1, // q=0
      -1.1, -2.1, -1.2, 2.2, // q=1
      -2.1, 0.1, 0.2, 1.2}; // q=2

    BEGIN_INSTRUMENTATION;
    for( int i0 = 0; i0 < m0; ++i0 )
        for( int j0 = 0; j0 < n0; ++j0 )
            for( int q0 = 0; q0 < (Q); ++q0 )
                for( int r0 = 0; r0 < (R); ++r0 )
                    y[i0*n0+j0] += weights[q0*(R)+r0] *
                        x[ ((q0+i0)%m0)*n0 + ((r0+j0)%n0) ];
    END_INSTRUMENTATION;
}

```

richard.m.veras@ou.edu

12

## Implementation

### Steps:

1. Instead of manually writing SSA-compliant code, we will instead:
  - a. Write a generator that takes in an operation description and a simple schedule, and
  - b. Automatically generate SSA-compliant C code as the output.
2. Implement the simple standalone code generator:
  - a. Read in an operation file.
  - b. Read an optional scheduling file.
  - c. Generate an SSA-compliant C code implementation.
  - d. Ensure SSA compliance.
  - e. Handles basic scheduling, if included.
  - f. Preserves the original operation logic.
3. Verify correctness:
  - a. Run `run-all-verify.sh`. (WIP)
4. Benchmark performance:
  - a. Run `run-all-measure.sh`. (WIP)
    - i. Generated SSA vs. Handwritten SSA
    - ii. Generated SSA vs. baseline (non-SSA)
5. Compare results:
  - a. Generated SSA vs. Handwritten SSA
  - b. Generated SSA vs. baseline (non-SSA)
  - c. Are there any discrepancies between manually written and generated SSA versions?
  - d. Does SSA that is generated produce unnecessary overhead?
6. Clean up generated files:
  - a. Run `run-clean-all.sh` to clean up all files generated by a given area. (WIP)

## T01: Simple Standalone + Unroll

Implement, Verify, Benchmark, and Compare for All of the Operation Specifications for All relevant schedule specifications in SSA form.

### Objective B: Standalone Simple (Unroll)

#### Operation Specification

```
3 4
-1.1 -1.1 1.2 -2.1
-1.1 -2.1 -1.2 2.2
-2.1 0.1 0.2 1.2
i0 j0 q0 r0
```

Standalone Code  
Generator

unroll r0

```
void COMPUTE_NAME( int m0,
int n0,
float *x,
float *y )
{
// BEGIN INSTRUMENTATION; // func: COMPUTE_NAME
for( int i0 = 0; i0 < m0; ++i0 )
// BEGIN INSTRUMENTATION; // loop:i0
for( int j0 = 0; j0 < n0; ++j0 )
// BEGIN INSTRUMENTATION; // loop:j0
for( int q0 = 0; q0 < (Q); ++q0 )
{
BEGIN_INSTRUMENTATION; // loop:q0
// r0=0
y[i0*n0+j0] += weights[q0*(R)+(0)] *
x[(q0*(R)+n0)*n0 + ((0)-j0)*n0];
}
// r0=1
y[i0*n0+j0] += weights[q0*(R)+(1)] *
x[(q0*(R)+n0)*n0 + ((1)-j0)*n0];
}
// r0=2
y[i0*n0+j0] += weights[q0*(R)+(2)] *
x[(q0*(R)+n0)*n0 + ((2)-j0)*n0];
}
// r0=3
y[i0*n0+j0] += weights[q0*(R)+(3)] *
x[(q0*(R)+n0)*n0 + ((3)-j0)*n0];
}
END_INSTRUMENTATION; // loop:q0
}
// END INSTRUMENTATION; // loop:j0
// END INSTRUMENTATION; // loop:i0
// END INSTRUMENTATION; // func:COMPUTE_NAME
}
```

richard.m.veras@ou.edu

13

### Implementation

Modify the code generator to support unrolling.

#### Steps:

1. Instead of manually writing SSA-compliant code, we will instead:
  - a. Write a generator that takes in an operation description and a simple unrolling schedule, and
  - b. Automatically apply unrolling to loops then generate SSA-compliant C code as the output.
2. Implement the simple standalone code generator:
  - a. Read in schedule file.
  - b. Generate an SSA-compliant C code implementation with modified loop structure according to schedule.
  - c. Ensure SSA compliance.
  - d. Preserves the original operation logic.
  - e. Store all unrolled SSA implementations in their own folder [i.e., /unrolled-versions/].
  - f. Ensure consistent naming across versions [i.e., `ssa\_unroll\_blur3x3.c`].
3. Verify correctness:
  - a. Run `run-all-verify.sh`. (WIP)
4. Benchmark performance:
  - a. Run `run-all-measure.sh`. (WIP)

- i. Unrolled SSA vs. non-unrolled SSA
  - ii. Unrolled SSA vs. baseline (non-SSA)
- 5. Compare results:
  - a. Unrolled SSA vs. non-unrolled SSA
  - b. Unrolled SSA vs. baseline (non-SSA)
- 6. Clean up generated files:
  - a. Run `run-clean-all.sh` to clean up all files generated by a given area. (WIP)

## T02: Simple Standalone + Split

Implement, Verify, Benchmark, and Compare for All of the Operation Specifications for All relevant schedule specifications in SSA form.

### Objective B: Standalone Simple (split)

Operation Specification

```
3 4
-1.1 -1.1 1.2 -2.1
-1.1 -2.1 -1.2 2.2
-2.1 0.1 0.2 1.2
i0 j0 q0 r0
```

Standalone Code  
Generator

split j0 j0\_o j0\_i 2

```
void COMPUTE_NAME( int m0,
                  int n0,
                  float *x,
                  float *y )
{
  for( int i0 = 0; i0 < m0; ++i0 )
  {
    for( int j0_o = 0; j0_o < n0; j0_o += 2 )
      for( int j0_i = 0; j0_i < 2; ++j0_i )
      {
        // BEGIN_INSTRUMENTATION; // loop:j0
        for( int q0 = 0; q0 < (0); ++q0 )
        {
          // BEGIN_INSTRUMENTATION; // loop:q0
          for( int r0 = 0; r0 < (R); ++r0 )
          {
            // BEGIN_INSTRUMENTATION; // loop:r0
            y[i0*n0+(j0_o+j0_i)] += weights[(i0*(R)+r0) *
              x[(q0*(R)+r0)*n0 + ((r0*(j0_o+j0_i))%n0) ];
            // END_INSTRUMENTATION; // loop:r0
          }
        }
      }
  }
}
```

richard.m.veras@ou.edu

14

## Implementation

Modify the code generator to support splitting.

### Steps:

1. Instead of manually writing SSA-compliant code, we will instead:
  - a. Write a generator that takes in an operation description and a simple splitting schedule, and
  - b. Automatically apply splitting to loops then generate SSA-compliant C code as the output.
2. Implement the simple standalone code generator:
  - a. Read in schedule file.
  - b. Generate an SSA-compliant C code implementation with modified loop structure according to schedule.
  - c. Ensure SSA compliance.

- d. Preserves the original operation logic.
- e. Store all split SSA implementations in their own folder [i.e., /split-versions/].
- f. Ensure consistent naming across versions [i.e., `ssa\_split\_blur3x3.c`].
3. Verify correctness:
  - a. Run `run-all-verify.sh`. (WIP)
4. Benchmark performance:
  - a. Run `run-all-measure.sh`. (WIP)
    - i. Split SSA vs. non-split SSA
    - ii. Split SSA vs. baseline (non-SSA)
5. Compare results:
  - a. Split SSA vs. non-split SSA
  - b. Split SSA vs. baseline (non-SSA)
6. Clean up generated files:
  - a. Run `run-clean-all.sh` to clean up all files generated by a given area. (WIP)

## T03: Simple Standalone + Interchange

Implement, Verify, Benchmark, and Compare for All of the Operation Specifications for All relevant schedule specifications in SSA form.

### Objective B: Standalone Simple (interchange)

Operation Specification

```
3 4
-1.1 -1.1 1.2 -2.1
-1.1 -2.1 -1.2 2.2
-2.1 0.1 0.2 1.2
i0 j0 q0 r0
```

Standalone Code  
Generator

Interchange r0 j0

```
void COMPUTE_NAME( int m0,
                  int n0,
                  float **,
                  float ** )
{
    // BEGIN INSTRUMENTATION; // func: COMPUTE_NAME
    for( int i0 = 0; i0 < m0; ++i0 )
    {
        // BEGIN INSTRUMENTATION; // loop:i0
        for( int r0 = 0; r0 < (R); ++r0 )
        {
            // BEGIN INSTRUMENTATION; // loop:r0
            for( int q0 = 0; q0 < (Q); ++q0 )
            {
                // BEGIN INSTRUMENTATION; // loop:q0
                for( int j0 = 0; j0 < n0; ++j0 )
                {
                    // BEGIN INSTRUMENTATION; // loop:j0
                    y[i0+r0+j0] += weights[i0*(R)+r0] *
                        x[(q0*(R)+q0)*n0 + ((r0+j0)%n0)];
                    // END INSTRUMENTATION; // loop:j0
                }
                // END INSTRUMENTATION; // loop:q0
            }
            // END INSTRUMENTATION; // loop:r0
        }
        // END INSTRUMENTATION; // loop:i0
    }
    // END INSTRUMENTATION; // func:COMPUTE_NAME
}
```

richard.m.veras@ou.edu

15

## Implementation

Modify the code generator to support loop interchanging.

### Steps:

1. Instead of manually writing SSA-compliant code, we will instead:
  - a. Write a generator that takes in an operation description and a simple interchanging schedule, and

- b. Automatically apply interchanging to loops then generate SSA-compliant C code as the output.
- 2. Implement the simple standalone code generator:
  - a. Read in schedule file.
  - b. Generate an SSA-compliant C code implementation with modified loop structure according to schedule.
  - c. Ensure SSA compliance.
  - d. Preserves the original operation logic.
  - e. Store all interchanged SSA implementations in their own folder [i.e., /interchanged-versions/].
  - f. Ensure consistent naming across versions [i.e., `ssa\_interchange\_blur3x3.c`].
- 3. Verify correctness:
  - a. Run `run-all-verify.sh`. (WIP)
- 4. Benchmark performance:
  - a. Run `run-all-measure.sh`. (WIP)
    - i. Interchanged SSA vs. non-interchanged SSA
    - ii. Interchanged SSA vs. baseline (non-SSA)
- 5. Compare results:
  - a. Interchanged SSA vs. non-interchanged SSA
  - b. Interchanged SSA vs. baseline (non-SSA)
- 6. Clean up generated files:
  - a. Run `run-clean-all.sh` to clean up all files generated by a given area. (WIP)

## Objective C: EoC Simple Schedule Code Generator

### T00: Implement AST to C codegen

Implement an abstract syntax tree (AST) to C code generator, and then test its correctness.

#### Implementation

### T01: Implement Operation Specification to AST

Implement and test.

### T02: Implement the Compiler Pass (AST to AST) for Implementing Unroll

Implement and test.

### T03: Implement the Compiler Pass (AST to AST) for Implementing Split

Implement and test.

## T04: Implement the Compiler Pass (AST to AST) for Implementing Interchange

Implement and test.

## T05: EoC Standalone Baseline

Implement, Verify, Benchmark, and Compare for All of the Operation Specifications in SSA form.

## T06: EoC Standalone + Unroll

Implement, Verify, Benchmark, and Compare for All of the Operation Specifications for All relevant schedule specifications in SSA form.

## T07: EoC Standalone + Split

Implement, Verify, Benchmark, and Compare for All of the Operation Specifications for All relevant schedule specifications in SSA form.

## T08: EoC Standalone + Interchange

Implement, Verify, Benchmark, and Compare for All of the Operation Specifications for All relevant schedule specifications in SSA form.

# Objective D: Complex Schedule Code Generator

## T00: Implement Generalized Schedule (AST to AST)

## T01: Complex Schedule Generator + Unroll + Unroll

## T02: Complex Schedule Generator + Split + Split + Split + Split

## T03: Complex Schedule Generator + Interchange + Interchange + Interchange + Interchange

## T04: Complex Schedule Generator + Mixed Schedule (length 3)

## T05: Complex Schedule Generator + Mixed Schedule (length 4)

## T06: Complex Schedule Generator + Mixed Schedule (length 5)

## Objective E: Sensitivity Analysis

T00: Analysis of Complex Schedule (length 3)

T01: Compare the Same Schedule Across Different Operations

T02: Sensitivity Study (vary the factor of a loop split)

T03: Sensitivity Study (vary architecture)

T04: Analysis of Complex Schedule (length 4)

T05: Compare the Same Schedule Across Different Operations

T06: Sensitivity Study (vary the factor of a loop split)

T07: Sensitivity Study (vary architecture)

T08: Analysis of Complex Schedule (length >4)

T09: Compare the Same Schedule Across Different Operations

T10: Sensitivity Study (vary the factor of a loop split)

T11: Sensitivity Study (vary the factor of a loop split for two split factors)

T12: Sensitivity Study (vary architecture)

T13: Post your best performance plot and schedule on Canvas

Optional: Repeats for Additional Points

T08-Opt00: Analysis of Complex Schedule (length >4)

T09-Opt00: Compare the Same Schedule Across Different Operations

T10-Opt00: Sensitivity Study (vary the factor of a loop split)

T11-Opt00: Sensitivity Study (vary the factor of a loop split for two split factors)

T12-Opt00: Sensitivity Study (vary architecture)

T13-Opt00: Post your best performance plot and schedule on Canvas



## Utilities & Testing